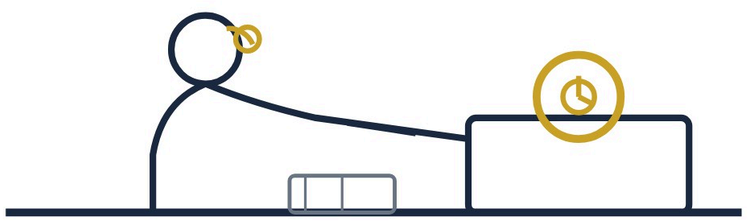
Product Development & MVP Strategy

Building Products People Actually Want

Three sessions of finance behind you.
Tonight we go back to building — carefully.

An MVP isn't a smaller product.
It's the smallest experiment that can prove you wrong.

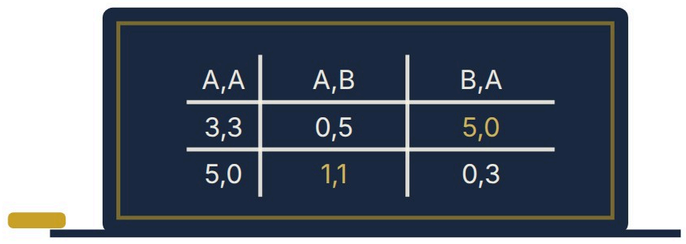
Two founders. Same idea. Same $60,000.



FOUNDER A

The Builder

Spends six months and the full $60K building the product she has in mind — complete, polished, ready for launch day.



FOUNDER B

The Scientist

Spends the same six months and the same $60K running four experiments — each one designed to kill a version of the idea.

VOTE · WHO IS FURTHER ALONG IN SIX MONTHS?

Founder A has a product — and **no idea whether anyone wants it.** Founder B has no product — and knows which three versions of the idea are dead, and which one somebody actually paid for. **One of them spent $60K on a guess. The other bought information.**

Time & Cash Burn = Startup Timebombs



Time

EVERY MONTH IS A MONTH OF RUNWAY

and a month your competitor, or the market, moves

&

Cash

EVERY DOLLAR SPENT BUILDING IS GONE

whether or not anyone wanted what you built

Every experiment you run spends both. The only question that matters:
what does it cost to find out — versus what does it cost to be wrong?

> Companies die from a lack of customers — not necessarily from a lack of cash.

The 4 Foundation Questions

1

What problem do you solve?

Named plainly — not your product, the problem.

2

Who's got it?

Groups you can describe, find, and call.

3

What's the impact today?

What it costs them right now — money, hours, risk.

4

Who pays to make it go away?

Not who likes it. Who *pays*.

Your answers to these are guesses. An MVP is a machine for testing them — one at a time, cheapest first. Question 4 is where tonight ends.

Four words people use interchangeably. They aren't.

Prototype	**Answers: can it work?** A test of feasibility, shown to yourselves. Nobody buys a prototype.
MVP	**Answers: will they use it — will they pay?** The smallest thing you can put in front of a real customer that produces a real decision. Tonight's word.
Pilot	**Answers: does it survive contact with one real customer?** Paid, scoped, delivered for real — the service business's natural MVP. More in Act II.
MLP	**Minimum Lovable Product — answers: will they tell anyone?** Enough polish that early users become your marketing. A later stage, not a starting point.

ACT II

Four companies you know.
What did Day One actually look like?

For each one: write down what you think they shipped first.
Then we'll look at what they actually did.

Zappos, 1999 — "people will buy shoes online."

Nick Swinmurn can't find the boots he wants at the mall. He believes people would buy shoes on the internet — in 1999, when everyone knows you'd never buy shoes you can't try on. **What does he build?**

> He builds almost nothing. He walks into local shoe stores, **photographs the inventory,** and posts the photos on a bare-bones site. When an order comes in, he goes back to the store, **buys the pair at full retail,** and ships it himself — losing money on every sale.



What it cost

a camera, a website, and shoe-store errands

bought

The answer

strangers will buy shoes they can't try on — proven with their own money

SOURCE

Swinmurn founded Zappos (as Shoesite.com) in 1999; sold to Amazon in 2009 for $1.2B. The store-photo test is the canonical "Wizard of Oz" MVP — manual labor behind a storefront that looks automated. Storefront shown is a recreation, not the original page.

DoorDash, 2013 — "restaurants need delivery."

Four Stanford students hear a macaroon shop owner say she has to turn away delivery requests. They believe restaurant delivery is a business. **What do they build?**

One afternoon of work: a landing page called **PaloAltoDelivery.com** — PDF menus found on the internet, and **their own cell number** at the bottom. They went back to their dorm rooms. The phone rang within hours: shrimp pad thai and egg rolls. They delivered it themselves.

The plan, if anyone called, was to say *"this is just a test — we're gathering data."* Then a real order came in from a real stranger, and they went and got the food. **The test converted them.**

SOURCE

Launched January 2013; first delivery January 12; $120K from Y Combinator that summer; renamed DoorDash. Founders did every delivery themselves for months — and DoorDash employees still do deliveries today.

Stripe, 2010 — "payments should be seven lines of code."

The Collison brothers believe developers hate their payment processors enough to switch. The product genuinely is a few lines of code. **So how do they get their first users?**

> When a founder said yes to trying the beta, they didn't send a link. They said **"Right then, give me your laptop"** — and installed it on the spot, watching every hesitation, every broken assumption, every question the user actually asked.

The link-senders learn whether people *say* they'll try it. The Collisons learned what happens when they actually *do* — which is a different answer entirely.

```
Stripe.setPublishableKey('pk_live_…');
Stripe.card.createToken({
  number: '4242 4242 4242 4242',
  exp_month: 8, exp_year: 2026,
  cvc: '123'
}, handleResponse);
// that's the whole integration.
```

SOURCE

Paul Graham named the move the "Collison installation" in his essay *Do Things That Don't Scale* — which is on tonight's resources slide. First user: Ross Boucher of 280 North, a YC friend. Code shown is illustrative of the era's API.

Franklin's Bagels — the $210,000 question.

Two weeks ago you priced this company. Year 3 of the model carries **$210K of capex for a second location** — and the whole valuation moved when it did. That number is a bet that a neighborhood Franklin's has never sold a bagel in will show up.

HOW DO YOU DE-RISK $210,000 FOR $2,000?



A folding table. **A farmers-market stall, six weekends, in the actual target neighborhood.** Same bagels, same prices. You learn foot traffic, repeat faces by week three, what sells out, and what a Saturday grosses — before anyone signs a lease.

The lease teaches you the same lesson for $210,000 — **after** you've signed it. The stall teaches it for $2,000, first. **No code anywhere in this MVP.**

The pattern hiding in all four stories.

Manual
BEATS AUTOMATED

Fake
BEATS REAL

Now
BEATS PERFECT

Every one of these bought a **customer decision** — not a product.

NOTICE THE SYMMETRY

The MVP of a **product company** is usually a **service** — a human photographing shoes, a founder delivering pad thai, a brother installing code by hand.

WHICH MEANS

The MVP of a **service business** is usually a **product** — a fixed-scope, fixed-price slice of what you'd otherwise custom-quote. Next slide.

What's the MVP when you sell hours?

MOVE 1

The pilot engagement

Paid, scoped, small, and honest about being a pilot. A real client paying a real invoice is worth fifty "we'd definitely buy that" conversations.

MOVE 2

The concierge test

Deliver by hand what you'd eventually systematize. Zappos in a blazer: the client sees a service; you see whether the process is worth building.

MOVE 3

Sell before you staff

A signed LOI for an offering that doesn't exist yet is the service world's preorder. If nobody signs, you just saved a hire.

All three test the same thing: **Question 4 — who's willing to pay?** Before you build the team, the system, or the office to deliver it.

Own the client relationship — or someone else will.



Vanguard

VALLEY FORGE, PA · EST. 1975 · TOLD FIRSTHAND

1975

Administration

Vanguard is born as the Wellington funds' **administrative arm** — the back office. Explicitly barred from distribution and investment advisory. The lowest-value rung on the ladder.

Then

Distribution

Within a few years it takes over **marketing and distribution** — and now it owns the shareholder relationship.

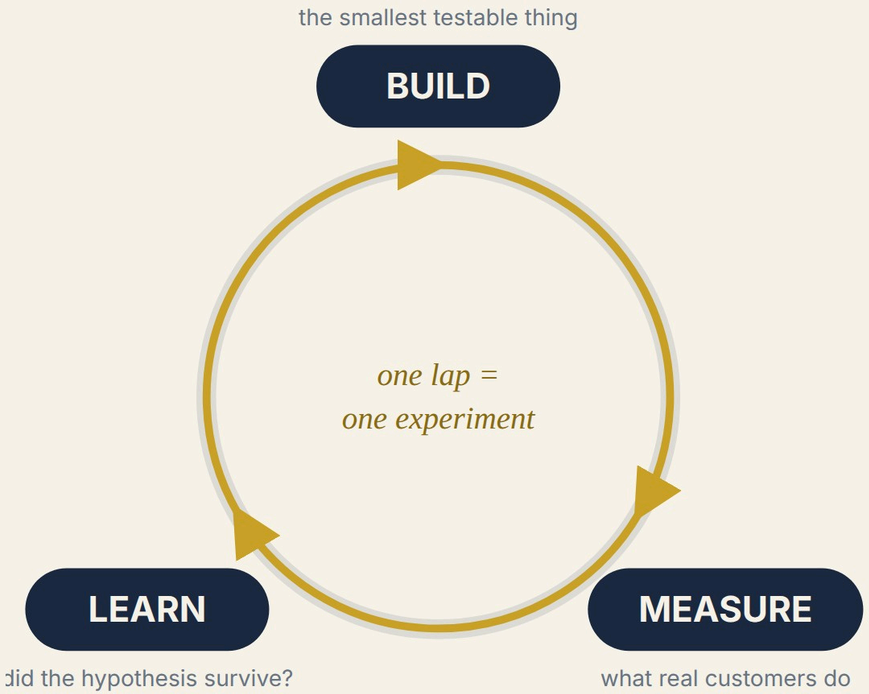
Then

Investment management

Then the **investment function itself.** The back office became the $10-trillion brand; the advisor became a vendor.

If you don't own the client relationship, you're a feature in someone else's product — and someone can run it on you. Every white-label shop and subcontractor in this room: this slide is for you.

The loop everyone quotes — and the part they skip.



The loop (Eric Ries, *The Lean Startup*): build the smallest thing that tests a hypothesis, measure what real customers do, learn whether the hypothesis survived — then go around again.

Measure the right thing. Vanity metrics go up and to the right and decide nothing: signups, page views, followers. Actionable metrics are behaviors that cost the customer something: **repeat use, referrals, money.**

The honest caveat: the loop assumes you already know what to measure — and which feature to test first. It doesn't tell you. That's the next fifteen minutes.

Carpenter Hall is packaging a new offering. What's essential?

Carpenter Hall Marketing is designing a monthly marketing package for small firms. Here's the candidate feature list. **Raise your hand for every feature you'd call essential.**

- Dedicated account manager
- Weekly reporting
- 48-hour turnaround
- Quarterly strategy workshop
- Monthly analytics review
- Social scheduling
- Ad management
- Content calendar

Everything got marked. It always does. Asked one at a time, every feature is essential — because saying yes costs nothing. A feature list rated feature-by-feature always comes back all-essential, and tells you exactly nothing about what to build first.

Stated importance is free.
Tradeoffs tell the truth.

YOU ALREADY KNOW THIS

S02: ask about **past behavior, not future intentions.** "Would you use this?" is worthless; "tell me about the last time" is evidence. Same rule, new domain.

THE FIX

Never ask what matters. **Force a choice between whole packages at real prices** — and read what they gave up. There's a sixty-year-old science for this — invented at Wharton, a mile up Walnut Street.

Conjoint analysis — invented at Wharton.

DEFINITION

Conjoint analysis is a survey method that measures what people actually value — by making them choose between complete packages instead of rating features one at a time.

The trick: it never asks what's important. It shows people complete packages — features bundled with a price — and makes them choose. Preference is inferred from **what they were willing to give up,** which is the only honest signal a hypothetical question can produce.

THE PEDIGREE

Developed by **Paul Green** in Wharton's marketing department in the mid-1960s; the landmark paper landed in 1971. It made him the most-cited author in marketing.

YOU'VE BEEN IN ONE

Conjoint designed **Courtyard by Marriott** and tested demand for **E-ZPass.** If you've driven the Schuylkill with a transponder, you're living inside a Wharton conjoint study.

SOURCE

Green & Rao, "Conjoint Measurement for Quantifying Judgmental Data," *Journal of Marketing Research*, 1971. Marriott and E-ZPass applications per Wharton Magazine and Knowledge at Wharton.

You run a small firm. Which package do you buy?

Package A	Package B	Package C	Package D
$2,500 PER MONTH	$4,000 PER MONTH	$4,000 PER MONTH	$5,500 PER MONTH
Project-based work	**Retainer** relationship	**Retainer** relationship	**Retainer** relationship
Shared support team	**Dedicated manager**	Shared support team	**Dedicated manager**
5-day turnaround	5-day turnaround	**48-hour turnaround**	**48-hour turnaround**
Monthly report	Monthly report	Monthly report	Monthly report

Hands up — one vote each. Notice what just happened to you: nobody asked whether a dedicated manager is important. You had to pay for it or give it up.

What conjoint won't tell you.

What it answers

What to build **first** — which features earn their place in version one, and which tier structure to test

≠

What it can't

Whether to build **at all.** A choice card is still a hypothetical — nobody's money moved

A customer who pays beats a customer who approves. Conjoint tells you what goes in the MVP — only a real purchase decision tells you whether the MVP lives.

And yes — putting price on those cards surfaced willingness-to-pay. That thread is Session 8. We're pulling it deliberately, not tonight.

Break — 10 minutes.

When we're back: we build a live product in this room, in real time —
and then ask why we shouldn't have.

Lean Startup rests on a premise. Check the date.



The Lean Startup is a 2011 book. Its core logic: **building software is slow and expensive** — months of engineering per attempt — so build the smallest thing that can teach you anything. The MVP was, at heart, **a cost-control device.**

2011

A LANDING PAGE + WORKING PROTOTYPE

weeks of engineering, tens of thousands of dollars

→

2026

THE SAME ARTIFACT

one prompt, a few minutes, effectively free

The premise collapsed. So does the discipline collapse with it?

Let's find out — live.

Carpenter Hall ships a product. Right now.

Carpenter Hall wants to test Move 3 from earlier — a productized, fixed-price offering, sold before it's staffed. Here's the actual prompt. Watch the clock.

```
> Build a single-file landing page for Carpenter Hall Marketing's new productized service: "The Marketing Audit — $2,500, delivered in two weeks." Fixed scope: channel review, competitor snapshot, 90-day action plan. One clear headline, three proof points, and an intake form (name, email, company, biggest marketing frustration). Professional, trustworthy, small-business tone. No frameworks, no dependencies.
```

While it builds, write down your prediction: **what will this page teach Carpenter Hall about the offering?**

This is the pattern, not a party trick: use AI to accelerate every step of MVP planning and execution — sharpening the hypothesis, drafting the test asset, writing the outreach. A frontier model does this best; tonight's build runs on **Claude Fable 5.** The Prompt Pack at the end of the deck chains the planning steps for you.

What we just did — and what we didn't.

WHAT WE HAVE

A real page. Real copy, a real form, a real price on a real offering. Ten minutes ago it didn't exist. Cost: one prompt and ten minutes.

WHAT WE KNOW

Nothing. Zero information about whether anyone wants a $2,500 audit. The page has taught us exactly as much as the unbuilt page: it hasn't met a customer yet.

The artifact was never the experiment. The experiment starts when a real prospect sees the page and decides. Everything before that moment is setup — and setup is now free. The customer's decision never will be.

"If building is free —
why bother with an MVP at all?
Just build the whole thing."

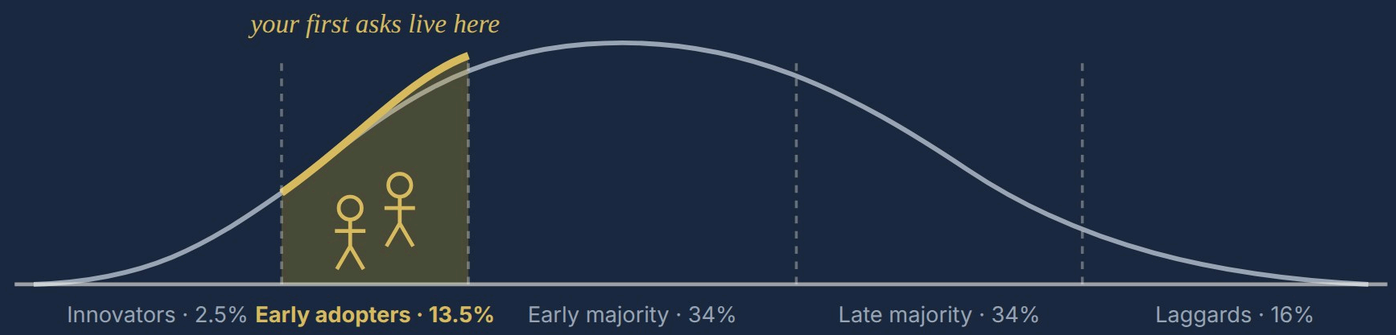
It's the right question, and the honest answer is: **the old reason is gone.** If the MVP was only ever about saving engineering cost, it's obsolete.

But the constraint didn't disappear. **It moved.**

The scarce resource was never engineering hours.

It's **customer attention.** You get a finite number of asks — emails answered, calls taken, pilots agreed to — before people stop replying. Spend them on a full product you guessed at, and you've burned the audience along with the build. **That's what an MVP conserves now.**

So spend your first asks where the signal is richest — **early adopters:**



TRAIT 1

Already in the problem

They're actively hacking together a solution today — which means they'll actually engage.

TRAIT 2

Unoffended by rough edges

Rough edges don't put them off — getting it first is the appeal.

TRAIT 3

Sought out by their circles

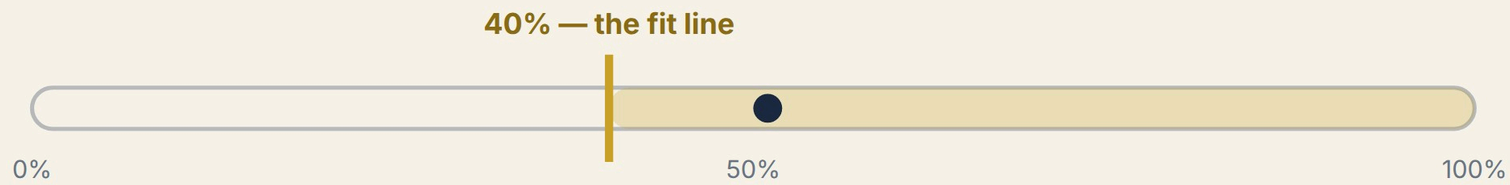
Known for expertise, asked for advice — they're where word of mouth starts.

The rest of the adoption curve — and the chasm between these people and everyone else — is Session 9.

How do you know it's working? One question.

"How would you feel if you could no longer use this product?"

Very disappointed | Somewhat disappointed | Not disappointed



Slack, 2015 open survey: 51%

40%+

"VERY DISAPPOINTED"

the product-market-fit signal — keep going, start scaling

< 40%

"VERY DISAPPOINTED"

keep iterating — growth spend on top of this leaks

SOURCE

Sean Ellis — first marketer at Dropbox, LogMeIn, and Eventbrite — derived the 40% threshold by comparing survey results across roughly 100 startups (published ~2009). The wording matters: disappointment, not satisfaction — loss reveals need; praise is free. Calibration point: Slack scored 51% in an open 2015 survey.

Meridian HR Advisory, month six. Fit — or not?

The setup: Meridian is a fictional 12-person HR consulting firm, six months after launching a monthly retainer product. Below is its actual dashboard — the same three instruments we just covered: a Sean Ellis score, pilot conversions, and retention curves. **Apply the 40% test and read the curves. You make the call.**

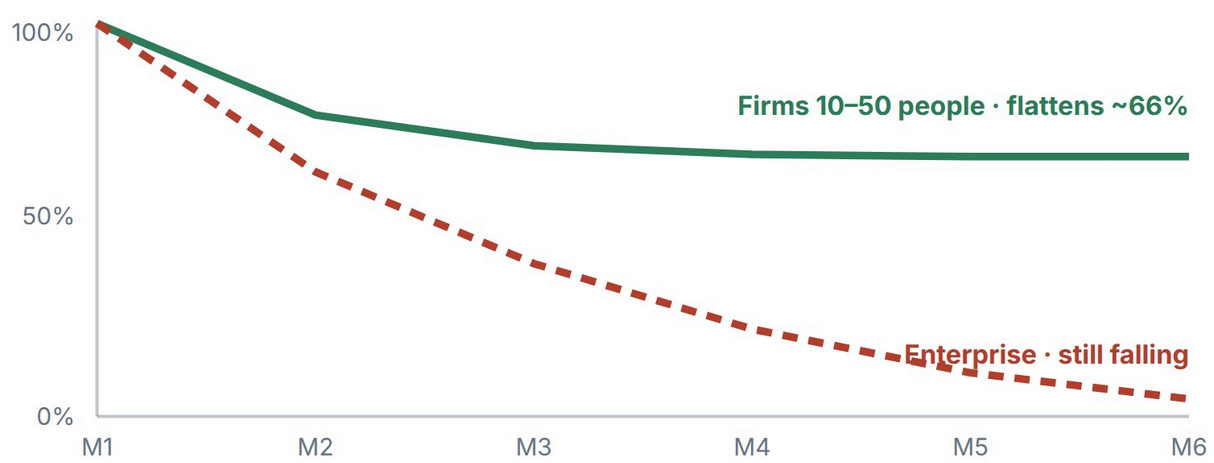
41%

"VERY DISAPPOINTED" · SEAN ELLIS SURVEY, N = 22 CLIENTS

9 of 20

PILOTS CONVERTED TO RETAINER THIS YEAR

MONTH-OVER-MONTH CLIENT RETENTION, BY SEGMENT



VOTE · DOES MERIDIAN HAVE PRODUCT-MARKET FIT?

THE LESSON

The honest read: yes — in one segment. The 41% clears the bar but n = 22 is thin. The tell is the shape of the curves: a retention curve that **flattens** is customers deciding to stay — that's fit. One that keeps **decaying** is a leaky bucket. Meridian has fit with 10–50-person firms and doesn't with enterprise. The move isn't "keep going" — it's **focus.**

Three signals that lie to you.

FALSE POSITIVE 1

Friends & warm intros

They're buying *you*, not the offering. Real signal starts with strangers — or at least friends' referrals, one step removed.

FALSE POSITIVE 2

Launch-day spikes

Novelty traffic decays by itself. Week-eight behavior is the truth; day-one behavior is a press release.

FALSE POSITIVE 3

Paid traffic

Money can simulate demand indefinitely. If growth stops when spend stops, you rented the signal — you don't have it.

And one honest sentence on pivots: a pivot isn't giving up — it's a change of strategy that **preserves what you learned.** Which is only possible if you ran real experiments. Founder A from slide 3 can't pivot; she can only start over.

Hypothesis Testing and Results Analysis

We believe **[customer]** will **[behavior]**.
We'll know within **[timeframe]** if **[measurable threshold]**.

FIX THIS ONE — OUT LOUD

"Clients want more strategic guidance."

Can this sentence lose? Then it can't win either. Rebuild it in the format above.

ONE REBUILD THAT CAN FAIL

"We believe **owners of 10–50-person firms** will **pay $500 for a quarterly strategy session**. We'll know within **3 weeks** if **3 of the next 10 clients we offer it to buy**."

One experiment at a time. Change two variables and you learn nothing about either.

A/B testing — honestly, at your scale.

THE PRINCIPLE

One variable at a time, threshold set ***before*** **you look.**

THE HONEST CAVEAT

Classic A/B is a traffic game — significance takes volume a thirty-client firm will never have. **Run the low-volume versions:**

LOW-VOLUME MOVE 1 · SEQUENTIAL TESTING

The next ten proposals go out as **version B** — new structure, new pricing page, new opener. Compare close rates against the last ten. Not lab-grade, but decision-grade.

LOW-VOLUME MOVE 2 · TEST WHAT YOU ALREADY SHIP

You don't have web traffic — you have **outreach emails, proposals, pricing pages.** Those are your funnel. **Your proposal template is your homepage.** Version it like one.

The discipline transfers even where the statistics can't:
decide what winning looks like before you peek.

The Experiment Card — five fields, no escape hatches.

1 — **Riskiest assumption**
The belief that, if wrong, kills the idea. Usually a Question 4 answer.

2 — **What result proves you wrong**
Written as a result you could actually observe — not a feeling.

3 — **Cheapest test that produces it**
Stall, pilot, page, LOI — the least you can do that forces a real decision.

4 — **Threshold — set before you run**
The number that means proceed, written down **before** the data arrives — so you can't move the goalposts after.

5 — **Deadline**
Experiments without deadlines are hobbies.

Look again — **it's the hypothesis format with a calendar attached.** The card link comes to you after tonight.

Carpenter Hall's card, filled in completely.

1. **Riskiest assumption:** small-firm owners will buy a fixed-price marketing audit **without a discovery call** — the productized offer sells itself from the page.

2. **Proves us wrong:** prospects read the page, then ask for a custom scope and a call anyway — or don't respond at all.

3. **Cheapest test:** tonight's landing page + the next **10 sales conversations** get pointed at it instead of a proposal.

4. **Threshold, set now: 3 of 10 buy** at $2,500 → build the offering for real. 1–2 → revise the offer. 0 → the assumption is dead; discovery calls are part of the product.

5. **Deadline:** September 1 — results on the table when we meet for Session 8.

The Prompt Pack — three interviews, chained.

P1

Foundation Interview

What it does: Claude interviews *you*, one question at a time, pushing each answer to the Four Foundation Questions until it names a specific customer, a specific problem, and what that problem costs today. Vague answers get challenged, not accepted.
What you walk away with: a one-page Foundation Brief — your four answers, sharpened, with the weakest one flagged as your riskiest assumption.

P2

Segment & Value Prop Explorer

What it does: paste in the P1 brief. Claude maps 3–5 candidate customer segments, writes the value proposition for each, and ranks them by which is cheapest to reach and test first.
What you walk away with: a segment map — a table that drops straight into your S03 Business Model Canvas, with a recommended first target.

P3

MVP Planner

What it does: paste in P1 + P2. Claude designs the cheapest test of your riskiest assumption — MVP format (stall, pilot, page, LOI), target early adopters, outreach plan, and first messaging.
What you walk away with: a one-page MVP plan, plus your Experiment Card pre-filled: assumption, disproof, test, threshold, deadline.

Each output feeds the next — **the whole framework, run end to end on your own business.** These are homework, not lecture. Links on the last slide.

An MVP isn't a smaller product. It's the smallest experiment that can prove you wrong.

Your assignment is one Experiment Card: one assumption, one test, one threshold set before you run it, one deadline — **two weeks from tonight.** Session 8 opens with your results.

RESOURCES

Let's get testing!

TONIGHT'S TOOLS

- **The Experiment Card** — link posted to the community after tonight
- **Prompt Pack P1–P3** — Foundation Interview · Segment Explorer · MVP Planner, posted alongside it
- **Carpenter Hall live-build page** — tonight's demo, warts and all

READING

- Paul Graham — *Do Things That Don't Scale* (the Collison installation essay)
- First Round Review — *The Path to Product-Market Fit*
- Eric Ries — *The Lean Startup* (read with tonight's caveats)
- Rob Fitzpatrick — *The Mom Test* (the S02 companion, still the best)

See you September 1 — bring a dead assumption and be proud of it.